

Multivariate Statistical Learning - A simple Anomaly Detection using K-means algorithm

Tommaso Puccetti
Studente presso Università degli studi di Firenze

September 23, 2019

Contents

1	Clustering: basics	2
1.1	Main steps	2
2	Non-hierarchical clustering: K-means	3
2.1	K-means: implementation	3
3	Anomaly detection using K-means	6
3.1	Implementation	7
3.1.1	Segmentation	7
3.1.2	Calculate centroids	8
3.1.3	Reconstruction	8
3.1.4	Detecting anomalies	11
3.1.5	Conclusion	12
3.1.6	Dataset	13

List of Figures

1	Example of K-means usage on artificial data.	5
2	Example of K-means usage on artificial data: output.	5
3	Noise dataset without anomaly.	6
4	Noise dataset with an anomlay.	7

5	Amazon's server CPU usage with a known anomaly	8
6	Centroids obtained using K-means algorithm	9
7	Noise signal.	10
8	68_95_99.7 rule.	11

1 Clustering: basics

Clustering analysis *is a collection of statistical methods that can be used to assign units to groups, where group members share some characteristic*. Possible usage are:

- **data reduction**: reduce data that are homogeneous (similar);
- **find natural clusters** and describe their unknown properties;
- **find** useful and suitable **groupings**;
- **find** unusual data objects (i.e. outlier detection).

The aim is to find a way of **grouping statistical units** (or variables) in a data set in such a way that: **units are similar within groups and dissimilar among groups**.

1.1 Main steps

The essentials steps for a clustering algorithm are the following:

1. **Define a metric** to evaluate distances between units;
2. **Define an algorithm** to form groups.

The grouping algorithm can be:

- Hierarchical;
- Non-hierarchical.

2 Non-hierarchical clustering: K-means

K-means clustering is one of the simplest and popular unsupervised machine learning algorithms. The aim is to find a partition of the dataset into a set of K clusters. To achieve that the algorithm takes a given K and then find a partition of K clusters that optimizes the chosen partitioning criterion. Here the main steps in details:

1. choose K points at random as cluster centers (centroids);
2. assign each unit to its closest centroid, using an opportune distance measure (usually Euclidean or Manhattan);
3. calculate the centroid of each cluster and use it as the new cluster center;
4. go back to step 2, stop when cluster centers do not change any more.

2.1 K-means: implementation

The implementation of the algorithm can be found in the *skmeans.py* file. The file contains **K_means** class declaration which implements the *fit()* method, used to find the centroids of the dataset, given K. The class also implements all the methods necessary to use K-means in the context of an anomaly detection algorithm. These methods are explained in details in the anomaly detection section.

1. First of all the algorithm choose K random centroid.

```
z = self.rng.permutation(len(X))[:self.k]
X = np.array(X, dtype= "float64")
centers = X[z]
```

2. The second step of the *fit()* method is to **calculate the Euclidean distance** between all the dataset points in respect to the K centroids. The distances are used to assign all the points to the nearest cluster, using K labels.

$$d(x, y) = \sqrt{\sum_{i=0}^n (x_i - y_i)^2}$$

```

for i in range(self.k):
    distances[:,i] = np.linalg.norm(X - centers[i], axis=1)

labels= np.argmin(distances, axis = 1)

```

3. For each cluster the algorithm **calculates the new centroid** as the mean of the points assigned to them:

```

for i in range(self.k):
    if(i in labels):
        new_centers = X[labels == i].mean(0)
        center_temp.append(new_centers)
    else:
        center_temp.append(self.farthest(X, distances, labels))

new_centers = np.array(center_temp, dtype= "float64")

```

This phase could generate an error in case one of the centroid has no segments assigned: in fact the computation of the mean for the centroid could not be done, as we try to divide for zero. To avoid that, the *farthest()* function is used to choose the farthest segment from the centroid, designating it as the new centroid.

```

def farthest(self, segments, distances, labels):
    counts = np.bincount(labels)
    new_centroid_idx = np.argmax(distances[:, np.argmax(counts)],
                                axis =0)
    new_centroid = segments[new_centroid_idx]
    print(new_centroid)

    return new_centroid

```

4. Repeats from point 2 to point 3 until the new centroids are equal to the centroids from the previous step.

```

if np.all(centers == new_centers):
    break

centers = new_centers

```

The *kmeans_sk.py* file contains an example of usage for the algorithm. The dataset presents 4 distinct points aggregates. Using the *fit()* method with $K = 4$ the algorithm, as we expected, identifies the 4 aggregates of points as 4 clusters.




Figure 1: Example of K-means usage on artificial data.




Figure 2: Example of K-means usage on artificial data: output.

3 Anomaly detection using K-means

Assume to have a record of how some values changes over time (i.e server usage, heartbeat data ecc), also assume this record having a regular shape or pattern that is indicative of some kind of "normal" behaviour. The K-means algorithm can be used to spot anomalies given the normal behaviour. The waveform in *Figure 3* rapresents a noise that is regular over time, instead, in *Figure 4*, an example of a possible anomaly is introduced.




Figure 3: Noise dataset without anomaly.

The idea is the following:

1. split the waveform into n segments of length l : every segment is considered a statistical unit;
2. apply the K-means algorithm over the l -dimension points obtained (segments). The results are K cluster's *centroids*;
3. a centroid is a segments itself that can be used to reconstruct the original waveform: for every segments the nearest centroid is used to approximate the current;
4. once obtained a recostruted waveform using only centroids, we can subtract it from the original: the result is also a waveform that rapresents the noise signal between the two;




Figure 4: Noise dataset with an anomaly.

5. the last step is to use a threshold value to detect where the error between the original and the reconstructed waveforms is anomalous in respect to the noise obtained.

3.1 Implementation

3.1.1 Segmentation

The first step is to define a training set. The training set is a waveform which represents the normal behaviour for our dataset. It is obvious that, the more we know about the normal behaviour, the more we are capable of spotting anomalies. The training set, once split into segments, will be used as input for the K-means algorithm. The segmentation is done using the *segmentation()* function implemented in the K-means class.

```
segment_len = 32
slide_len = 16
km = skmeans.K_means(30)
segments = km.segmentation(train_array, segment_len, slide_len)
```

The idea is to have a sliding window which is translated to obtain overlapping segments. This function simply takes the training set as input, along with the desired length for the segments and the slide length for the window. The slide length is chosen to be half of the segments length, in this way

the reconstruction could be easily done in further steps. The output would be a Numpy's array containing n segments. The dataset, shown in picture from now on, refers to Amazon server CPU usage with a known anomaly.




Figure 5: Amazon's server CPU usage with a known anomaly

3.1.2 Calculate centroids

The *fit()* function of the K-means class, described in the previous chapter, is used to calculate the centroid given the segments obtained from the previous step.

```
centr = km.fit(segments)
learn_utils.plot_waves(centr, step=3)
```

The output would be a Numpy's array containing K centroids of the dataset. The picture below shows some of the normal waveform obtained.

3.1.3 Reconstruction

Once obtained the centroids we can use them to reconstruct the original waveform. First of all we have to apply again the segmentation process, this time over the entire dataset. Then the *predict()* function is used to find, for all the segments obtained, the nearest centroid. The output would be an array which store in the n-th position a label referring to the nearest centroid for the n-th segment.

```
segments = km.segmentation(dataset_array, segment_len, slide_len)
lab = km.predict(centr, segments)
```





Figure 6: Centroids obtained using K-means algorithm

```
def predict(self, centers, windowed_segments):
    distances = np.zeros((len(windowed_segments), self.k))
    for i in range(self.k):
        distances[:,i] = np.linalg.norm(windowed_segments
                                         - centers[i], axis=1)

    labels= np.argmin(distances, axis = 1)
    print("Predicted Labels")
    print(labels)

    return labels
```

```
Predicted Labels
[12 27  3 14 23 22 27 16 12 26 12 19  4  5  7  6  6  6 28 10  9  1 10  4
29 16  6 21 20 15 13  9 24  7 12  2 17 17 11  1 23  9 10  9 11 11 10  6
 6  8  0  8  9  9  1 11  0 23  0 21 26 27  6  8 24 23 25 25 10 23 25  8
 9 10 25  1 11 10  5 23 24  0 10  8  9 10 23 23  0 23 23 18 10 10  7 23
 6  5 23  7 21 26 21 15 13 27 14 27 14 21 15 15 15 15 15 15 13 13 21
20 12 15 26 26 20 15 15 26 15 26 15 15 21 26 21 26 27 27 12 27 27 23 27
27 21 26 27 21 21 20 15 26 21 26 26 15 27 15 21 26 26 21 15 21 21 26 26
21 26 26 15 15 26 21 21 26 21 21 26 21 21 26 21 15 26 21 15 26 21 15 21
20 21 26 21 26 26 21 21 26 21 21 26 20 27  6  7 27 12 27 14  6 21 21 15
21 13 13 27  6 10  0 23  9  0 12 21 12 12 20 20 27 12 27 12 27 21 21 20
23 23  9  0  9 10 10  8  9]
```

Finally the reconstruction algorithm is used to obtain the whole waveform.

```
noise = detection.reconstruction(dataset_array, segments, lab,
                                centr, slide_len, segment_len)
```

```

def reconstruction(data, segments, label, centroid, slide_len, segment_len):
    n = 0
    reconstruction = np.zeros(len(data))
    win_seg = np.zeros(len(data))
    pos = 0
    for segment in segments:
        nearest_centroid = centroid[label[n]]
        win_seg[pos:pos+segment_len] += segment
        reconstruction[pos:pos+segment_len] += nearest_centroid
        pos = n * slide_len
        n = n + 1

    noise = reconstruction[0:len(reconstruction)] -
        win_seg[0:len(win_seg)]
    n_plot_samples = len(noise)

    plt.plot(win_seg[0:n_plot_samples], label="Original waveform")
    plt.plot(reconstruction[0:n_plot_samples],
             label="Reconstructed waveform")
    plt.plot(noise[0:n_plot_samples], label="Noise")
    plt.legend()
    plt.show()
    plt.hist(noise)
    plt.show()

    return noise, np.mean(noise)

```

The noise obtained as output is plotted along with the original and the reconstructed waveform.




Figure 7: Noise signal.

3.1.4 Detecting anomalies

The noise obtained from previous step can be analyzed to detect anomalies. To do that we defined the noise signal as being normally distributed with a mean of zero. If it suddenly changes into another distribution, something unusual or anomalous is happening. To spot outliers values in the noise distribution we use a threshold value defined as:

$$mean \pm 4 * std$$

Based on the normal distribution properties, values lower than $mean - 4 * std$ deviation and higher than $mean + 4 * std$ should be extremely rare and could be considered as an anomaly:

- 68% of all values fall between $[mean - \sigma; mean + \sigma]$;
- 95% of all values fall between $[mean - 2 * \sigma; mean + 2 * \sigma]$;
- 99,7% of all values fall between $[mean - 3 * \sigma; mean + 3 * \sigma]$

These 3 rules are also known as the 68_95_99.7 rule.




Figure 8: 68_95_99.7 rule.

```

pos_treshold = np.mean(noise) + coef * np.std(noise)
neg_treshold = np.mean(noise) - coef * np.std(noise)
anomalies = []

for i in range(0, len(noise)):
    if(noise[i] <= neg_treshold):
        anomalies.append((i, noise[i]))
    if(noise[i] > pos_treshold):
        anomalies.append((i, noise[i]))

```

The *if* branches are used to detect and store anomalies in a specific Python list, returned as output of the function.

```

[[1752.    136.241]
 [1753.    136.351]
 [1754.    137.356]
 [1755.    137.524]
 [1756.    131.364]
 [1757.    132.855]
 [1758.    141.735]
 [1759.    106.452]
 [1760.    132.65 ]
 [1761.    134.158]
 .
 .
 .
 [1875.    133.178]
 [1876.    127.938]
 [1877.    135.216]
 [1878.    135.596]
 [1879.    109.638]
 [1880.    136.132]]

```

3.1.5 Conclusion

The first part of the code (noise production) can be used as base to implement an anomaly detection with use different approach for detection. For example, a sliding window can be used to detect changes in the noise distribution rather than use a threshold value to detect outliers values. The idea is to use window containing the *n* most recent elements. If the mean and standard derivation of this window change too much from their expected values, we can deduce an anomaly. The use of the sliding windows could be combined with the outliers detection to improve the detection precision.



The code of the project can be reached at: <https://github.com/GeorgeSmiley/masl>.

3.1.6 Dataset

The dataset used is the Numenta Anomaly Benchmark (NAB) that is a benchmark for evaluating algorithms for anomaly detection in streaming and online applications. It is comprised of over 50 labeled real-world and artificial timeseries data files plus a novel scoring mechanism designed for real-time applications. All of the data and code is fully open-source, with extensive documentation, and a scoreboard of anomaly detection algorithms: github.com/numenta/NAB.